

The key here is that, while this kind of connectivity includes a notion of time, it can be converted into a graph connectivity problem of a more standard sort.

We construct the following directed graph G . We scan through the ordered triples in the trace data, maintaining an array pointing to linked lists associated with each computer C_a . (Each list is initialized to null.)

For each triple (C_i, C_j, t_k) we see in our scan, we create nodes (C_i, t_k) and (C_j, t_k) , and we create directed edges joining these two nodes in both directions. Also, we append these nodes to the lists for C_i and C_j respectively. If this is not the first triple involving C_i , then we include a directed edge from (C_i, t) to (C_i, t_k) , where t is the timestamp in the preceding element (the previously last one) in the list for C_i . We do the analogous thing for C_j . By explicitly maintaining these lists for each node, we are thus able to construct all these new nodes and edges in constant time per triple.

Now, given a collection of triples, we want to decide whether a virus introduced at computer C_a at time x could have infected computer C_b by time y . We walk through the list for C_a until we get to the last node (C_a, x') for which $x' \leq x$. We now run directed BFS from (C_a, x') to determine all nodes that are reachable from it. If a node of the form (C_b, y') with $y' \leq y$ is reachable, then we declare that C_b could have been infected by time y ; otherwise we declare it could not have.

Let's argue first about the correctness of the algorithm, then its running time. First, we claim that if there is a path from (C_a, x') to (C_b, y') as in the previous paragraph, then C_b could have been infected by time y . To see this, we simply have the virus move between computers C_i and C_j at time t_k , whenever an edge from (C_i, t_k) to (C_j, t_k) is traversed by the BFS. This is a feasible sequence of virus transmissions that results in the virus first leaving C_a at time x or later (by the definition of x') and arriving at C_b by time y .

Conversely, suppose there were a sequence of virus transmissions that results in the virus first leaving C_a at time x or later and arriving at C_b by time y . Then we can build a path in our graph as follows. We start at node (C_a, x') and follow edges to (C_a, x'') , for the x'' when the virus first leaves C_a . (Note that there are such edges since $x' \leq x \leq x''$; or else $x' = x''$.) In general, for each time that the virus moves from C_i to C_j at time t_k , we add the edge from (C_i, t_k) to (C_j, t_k) to the path; if it next moves out of C_j at time $t \geq t_k$, we add the the sequence of edges from (C_j, t_k) to (C_j, t) . When the virus first arrives at node C_b , we will have just added a node (C_b, y'') to the path; since $y'' \leq y$ and y' is the largest y involving C_b in the trace data with this property, there is a sequence of edges from (C_b, y'') to (C_b, y') , completing the path.

Finally, we consider the running time. Each triple in the trace data causes us to add a constant number of nodes and edges to the graph, so the graph has $O(m)$ nodes and edges, and since we build it in constant time per node and edge, this takes time $O(m)$. Running BFS takes time linear in the size of the graph, so this too takes time $O(m)$.

¹ex207.316.912